

## class09

Jibin (PID: A53300326)

2021/10/27

### Exploratory data analysis

Use the `read.csv()` function to read the CSV (comma-separated values) file containing the data (available from our class website: `WisconsinCancer.csv` )

```
fna.data <- read.csv("https://bioboot.github.io/bimm143_S20/class-material/WisconsinCancer.csv")
```

Complete the following code to input the data and store as `wisc.df`

```
wisc.df <- data.frame (fna.data, row.names=1)
```

We can use -1 here to remove the first column

```
wisc.data <- wisc.df[,-1]
```

Create diagnosis vector for later

```
diagnosis <- factor(wisc.df$diagnosis)
diagnosis
```

##	[1]	M M M M M M M M M M M M M M M B B B M M M M M M M M M M M M M
##	[38]	B M M M M M M M M B M B B B B M M B M M B B B B M B M M B B B B M M M
##	[75]	B M B M M B B B M M B M M M B B B M B B M M B B B M M B B B B M B B M B B
##	[112]	B B B B B B M M M B M M B B B M M B M B M M B M M B B M B B B M B B B M B
##	[149]	B B B B B B B M B B B B M M B M B B M M B B M M B B B B M B B M M M B M
##	[186]	B M B B B M B B M M B M M M M B M M M B M B M B B M B M M M M B B M M B B
##	[223]	B M B B B B B M M B B M B B M M B M B B B B M B B B B M B M M M M M M M
##	[260]	M M M M M M M B B B B B B M B M B B B M B B M M B B B B B B B B B B B B
##	[297]	B M B B M B M B B B B B B B B B B B B B B M B B B M B M B B B B M M M B B
##	[334]	B B M B M B M B B B M B B B B B B B B M M M B B B B B B B B B B M M B M M
##	[371]	M B M M B B B B B M B B B B B M B B B M B B M M B B B B B B M B B B B B B
##	[408]	B M B B B B B M B B M B B B B B B B B B B B B M B M M B M B B B B B M B B
##	[445]	M B M B B M B M B B B B B B B B M M B B B B B B M B B B B B B B B B B M B

```
## [482] B B B B B M B M B B M B B B B M M B M B M B B B B M B B M B M B M
## [519] B B B M B B B B B B B B B B M B M M B B B B B B B B B B B B B B
## [556] B B B B B B M M M M M M M B
## Levels: B M
```

Q1. How many observations are in this dataset?

```
dim(wisc.data)
```

```
## [1] 569 30
```

Q2. How many of the observations have a malignant diagnosis?

```
table(diagnosis)
```

```
## diagnosis
##      B      M
## 357 212
```

Q3. How many variables/features in the data are suffixed with `_mean`?

```
grep("_mean", colnames(wisc.data))
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

## *Principal Component Analysis*

### Performing PCA

Check column means and standard deviations

```
colMeans(wisc.data)
```

```
##           radius_mean      texture_mean      perimeter_mean
##      1.412729e+01      1.928965e+01      9.196903e+01
##           area_mean      smoothness_mean      compactness_mean
##      6.548891e+02      9.636028e-02      1.043410e-01
##      concavity_mean      concave.points_mean      symmetry_mean
##      8.879932e-02      4.891915e-02      1.811619e-01
## fractal_dimension_mean      radius_se      texture_se
##      6.279761e-02      4.051721e-01      1.216853e+00
##           perimeter_se      area_se      smoothness_se
##      2.866059e+00      4.033708e+01      7.040979e-03
##      compactness_se      concavity_se      concave.points_se
##      2.547814e-02      3.189372e-02      1.179614e-02
##           symmetry_se      fractal_dimension_se      radius_worst
##      2.054230e-02      3.794904e-03      1.626919e+01
##      texture_worst      perimeter_worst      area_worst
```

```
##          2.567722e+01          1.072612e+02          8.805831e+02
##      smoothness_worst      compactness_worst      concavity_worst
##          1.323686e-01          2.542650e-01          2.721885e-01
##      concave.points_worst      symmetry_worst fractal_dimension_worst
##          1.146062e-01          2.900756e-01          8.394582e-02
```

```
apply(wisc.data,2,sd)
```

```
##          radius_mean          texture_mean          perimeter_mean
##          3.524049e+00          4.301036e+00          2.429898e+01
##          area_mean          smoothness_mean          compactness_mean
##          3.519141e+02          1.406413e-02          5.281276e-02
##          concavity_mean      concave.points_mean          symmetry_mean
##          7.971981e-02          3.880284e-02          2.741428e-02
##      fractal_dimension_mean          radius_se          texture_se
##          7.060363e-03          2.773127e-01          5.516484e-01
##          perimeter_se          area_se          smoothness_se
##          2.021855e+00          4.549101e+01          3.002518e-03
##          compactness_se          concavity_se          concave.points_se
##          1.790818e-02          3.018606e-02          6.170285e-03
##          symmetry_se      fractal_dimension_se          radius_worst
##          8.266372e-03          2.646071e-03          4.833242e+00
##          texture_worst          perimeter_worst          area_worst
##          6.146258e+00          3.360254e+01          5.693570e+02
##          smoothness_worst      compactness_worst          concavity_worst
##          2.283243e-02          1.573365e-01          2.086243e-01
##      concave.points_worst      symmetry_worst fractal_dimension_worst
##          6.573234e-02          6.186747e-02          1.806127e-02
```

Perform PCA on wisc.data by completing the following code

```
wisc.pr <- prcomp(wisc.data, scale=TRUE)
summary(wisc.pr)
```

```
## Importance of components:
##          PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  3.6444 2.3857 1.67867 1.40735 1.28403 1.09880 0.82172
## Proportion of Variance 0.4427 0.1897 0.09393 0.06602 0.05496 0.04025 0.02251
## Cumulative Proportion 0.4427 0.6324 0.72636 0.79239 0.84734 0.88759 0.91010
##          PC8      PC9      PC10     PC11     PC12     PC13     PC14
## Standard deviation  0.69037 0.6457 0.59219 0.5421 0.51104 0.49128 0.39624
## Proportion of Variance 0.01589 0.0139 0.01169 0.0098 0.00871 0.00805 0.00523
## Cumulative Proportion 0.92598 0.9399 0.95157 0.9614 0.97007 0.97812 0.98335
##          PC15     PC16     PC17     PC18     PC19     PC20     PC21
## Standard deviation  0.30681 0.28260 0.24372 0.22939 0.22244 0.17652 0.1731
## Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
## Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
##          PC22     PC23     PC24     PC25     PC26     PC27     PC28
## Standard deviation  0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
## Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
## Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
```

##		PC29	PC30
##	Standard deviation	0.02736	0.01153
##	Proportion of Variance	0.00002	0.00000
##	Cumulative Proportion	1.00000	1.00000

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

**44.27%**

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

**3 PCs are required to describe at least 70% of the original variance in the data.**


Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

**7 PCs are to describe at least 90% of the original variance in the data.**

## Interpreting PCA results

Create a biplot of the `wisc.pr` using the `biplot()` function.

```
biplot(wisc.pr)
```




Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?


This is a hot mess of a plot and it's difficult to interpret.

Scatter plot observations by components 1 and 2

```
plot(wisc.pr$x[,1:2], xlab = "PC1", ylab = "PC2", col=diagnosis)
```



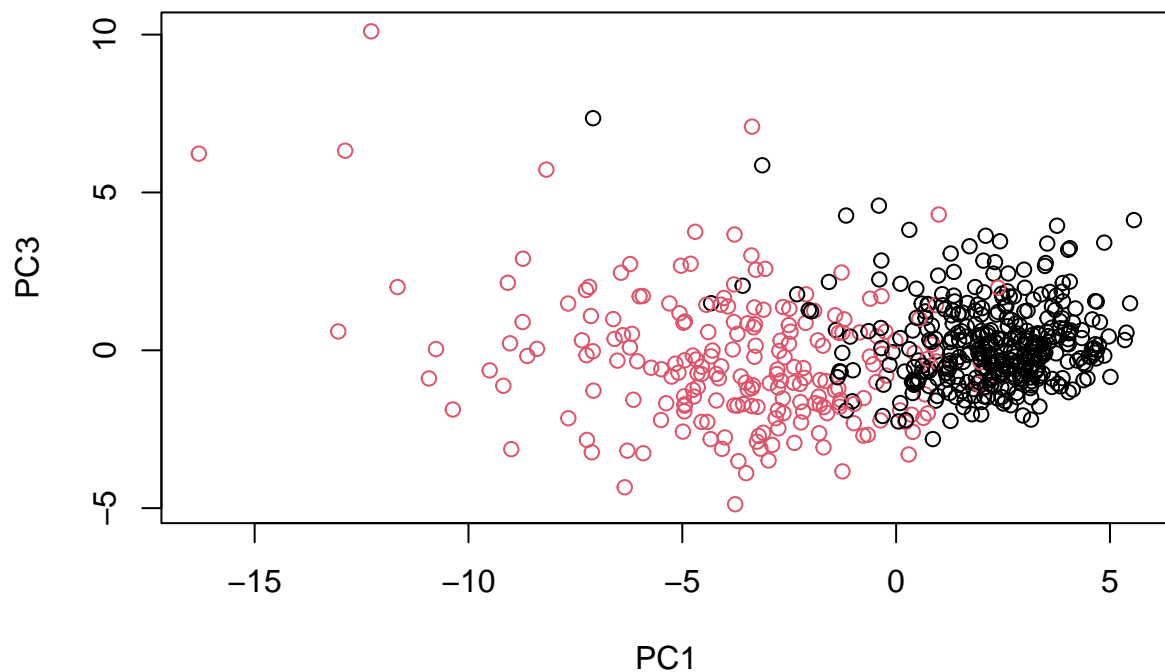
```
plot(wisc.pr$x [,1], wisc.pr$x [,3])
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

**Repeat for components 1 and 3**

```
plot(wisc.pr$x[, 1], wisc.pr$x[, 3], col = diagnosis,  
     xlab = "PC1", ylab = "PC3")
```



The first plot (PC2 vs PC1) has a cleaner cut separating the two subgroups than this plot, since PC 2 explains more variance in the original data than PC 3.

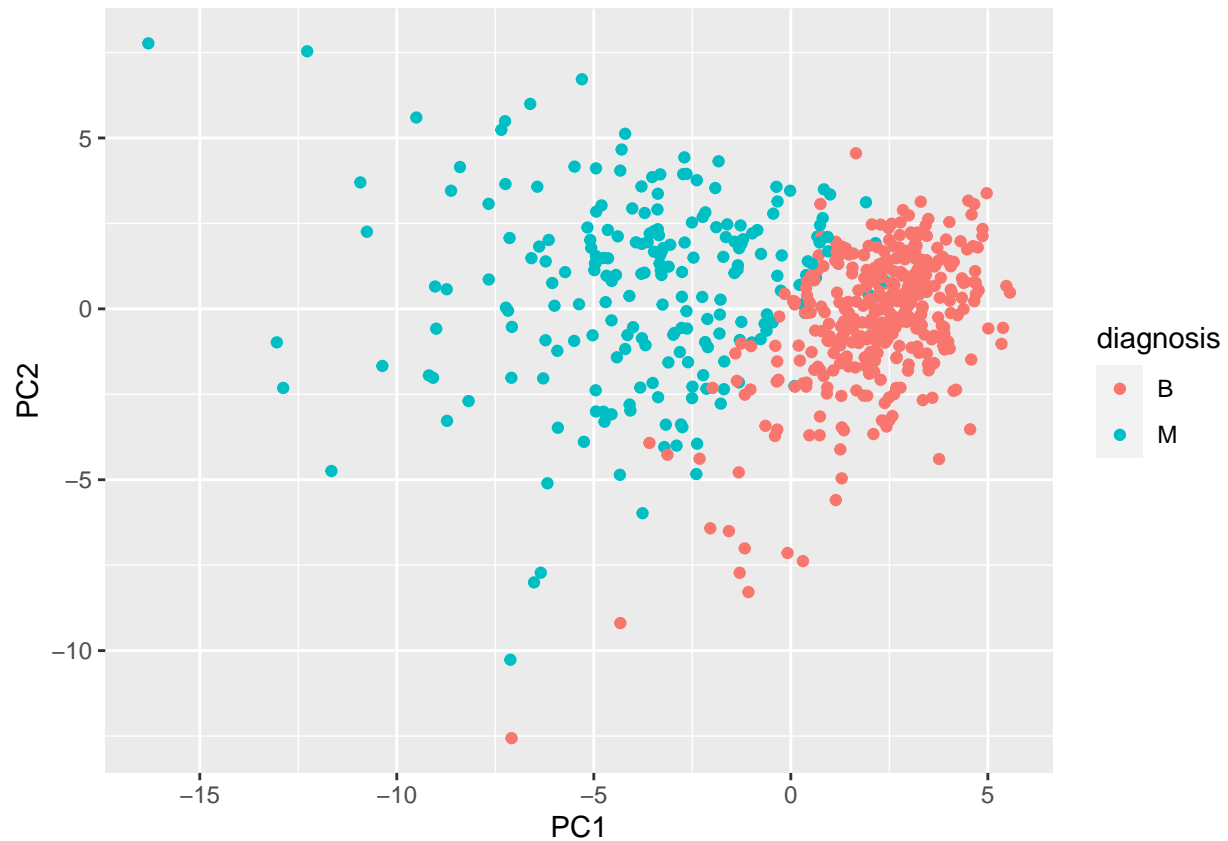
Create a data.frame for ggplot

```
df <- as.data.frame(wisc.pr$x)
df$diagnosis <- diagnosis
```

Make a scatter plot colored by diagnosis

```
library("ggplot2")
ggplot(df) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```





## Variance explained

Calculate variance of each component

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
## [1] 13.281608  5.691355  2.817949  1.980640  1.648731  1.207357
```


Variance explained by each principal component: pve

```
pve <- pr.var / sum(pr.var)
pve
```

```
## [1] 4.427203e-01 1.897118e-01 9.393163e-02 6.602135e-02 5.495768e-02
## [6] 4.024522e-02 2.250734e-02 1.588724e-02 1.389649e-02 1.168978e-02
## [11] 9.797190e-03 8.705379e-03 8.045250e-03 5.233657e-03 3.137832e-03
## [16] 2.662093e-03 1.979968e-03 1.753959e-03 1.649253e-03 1.038647e-03
## [21] 9.990965e-04 9.146468e-04 8.113613e-04 6.018336e-04 5.160424e-04
## [26] 2.725880e-04 2.300155e-04 5.297793e-05 2.496010e-05 4.434827e-06
```


Plot variance explained for each principal component

```
plot(pve, xlab = "Principal Component",  
     ylab = "Proportion of Variance Explained",  
     ylim = c(0, 1), type = "o")
```



Alternative scree plot of the same data, note data driven y-axis


```
barplot(pve, ylab = "Precent of Variance Explained",  
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)  
axis(2, at=pve, labels=round(pve,2)*100 )
```



ggplot based graph

```
#install.packages("factoextra")
```


```
library(factoextra)
fviz_eig(wisc.pr, addlabels = TRUE)
```



## Communicating PCA results

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`?

```
barplot(wisc.pr$rotation[,1], las=2, mar=c(10, 3, 3, 20))
```



### concave.points\_mean

Q10. What is the minimum number of principal components required to explain 80% of the variance of the data

```
var <- summary(wisc.pr)
var$importance
```

```
##          PC1      PC2      PC3      PC4      PC5      PC6
## Standard deviation  3.644394 2.385656 1.678675 1.407352 1.284029 1.098798
## Proportion of Variance 0.442720 0.189710 0.093930 0.066020 0.054960 0.040250
## Cumulative Proportion 0.442720 0.632430 0.726360 0.792390 0.847340 0.887590
##          PC7      PC8      PC9      PC10     PC11
## Standard deviation  0.8217178 0.6903746 0.6456739 0.5921938 0.5421399
## Proportion of Variance 0.0225100 0.0158900 0.0139000 0.0116900 0.0098000
## Cumulative Proportion 0.9101000 0.9259800 0.9398800 0.9515700 0.9613700
##          PC12     PC13     PC14     PC15     PC16
## Standard deviation 0.5110395 0.4912815 0.3962445 0.3068142 0.2826001
## Proportion of Variance 0.0087100 0.0080500 0.0052300 0.0031400 0.0026600
## Cumulative Proportion 0.9700700 0.9781200 0.9833500 0.9864900 0.9891500
##          PC17     PC18     PC19     PC20     PC21
## Standard deviation 0.2437192 0.2293878 0.2224356 0.1765203 0.1731268
## Proportion of Variance 0.0019800 0.0017500 0.0016500 0.0010400 0.0010000
## Cumulative Proportion 0.9911300 0.9928800 0.9945300 0.9955700 0.9965700
##          PC22     PC23     PC24     PC25     PC26
## Standard deviation 0.1656484 0.1560155 0.1343689 0.1244238 0.0904303
```

```
## Proportion of Variance 0.0009100 0.0008100 0.0006000 0.0005200 0.0002700
## Cumulative Proportion 0.9974900 0.9983000 0.9989000 0.9994200 0.9996900
##               PC27      PC28      PC29      PC30
## Standard deviation 0.08306903 0.0398665 0.02736427 0.01153451
## Proportion of Variance 0.00023000 0.0000500 0.00002000 0.00000000
## Cumulative Proportion 0.99992000 0.9999700 1.00000000 1.00000000
```

```
sum(var$importance[3,] <0.8)
```

```
## [1] 4
```

So, four at least.

## Hierarchical clustering

Scale the wisc.data data using the “scale()” function???


```
data.scaled <- scale(wisc.data)
data.dist <- dist(data.scaled)
wisc.hclust <- hclust(data.dist)
```

Results of hierarchical clustering

Q11. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)
abline(wisc.hclust, h=18, col="red", lty=2)
```

## Cluster Dendrogram



```
data.dist
hclust(*, "complete")
```

### Selecting number of clusters

Use `cutree()` to cut the tree so that it has 4 clusters. Assign the output to the variable `wisc.hclust.clusters`.

```
wisc.hclust.clusters <- cutree(wisc.hclust, 4)
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##                   1 12 165
##                   2  2   5
##                   3 343  40
##                   4  0   2
```

Q12. Can you find a better cluster vs diagnoses match by cutting into a different number of clusters between 2 and 10?

Perhaps no.


```
wisc.hclust.clusters <- cutree(wisc.hclust, 5)
table(wisc.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.hclust.clusters  B  M
##              1  12 165
##              2   0   5
##              3 343  40
##              4   2   0
##              5   0   2
```

Q13. Which method gives your favorite results for the same data.dist dataset? Explain your reasoning.

The method “ward.D2” looked much better, as the two major clusters are more clearly separated in this method

```
plot(hclust(data.dist, method = "average"))
```

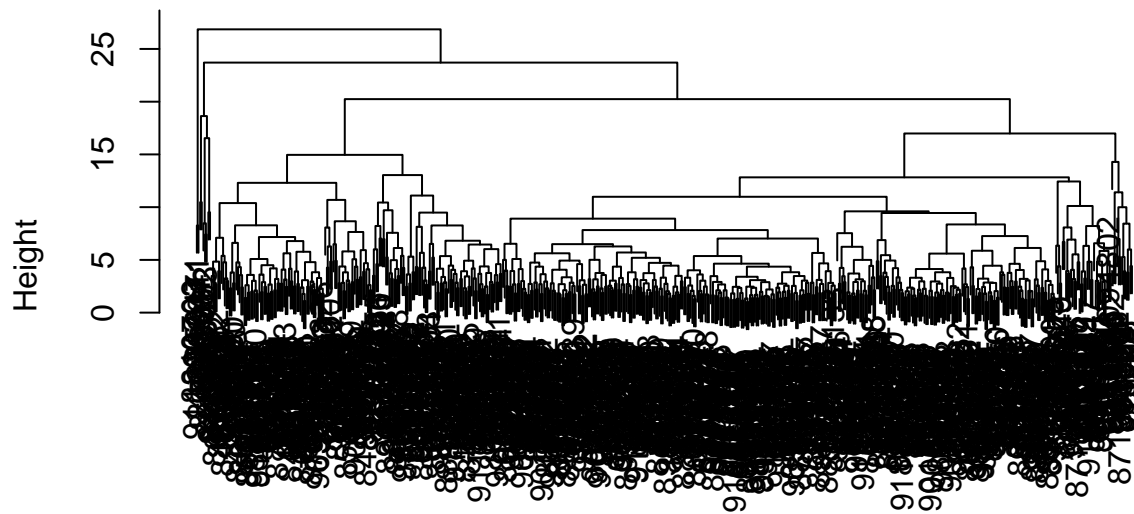


```
data.dist
hclust (*, "average")
```

```
plot(hclust(data.dist, method = "complete"))
```




## Cluster Dendrogram



data.dist  
hclust (\*, "complete")

```
plot(hclust(data.dist, method = "single"))
```


## Cluster Dendrogram



```
data.dist  
hclust (*, "single")
```

```
plot(hclust(data.dist, method = "ward.D2"))
```

## Cluster Dendrogram



```
data.dist
hclust (*, "ward.D2")
```

## OPTIONAL: K-means clustering

### K-means clustering and comparing results

```
wisc.km <- kmeans(wisc.data, centers= 2, nstart= 20)
table(wisc.km$cluster, diagnosis )
```

```
##      diagnosis
##      B      M
## 1 356    82
## 2   1   130
```

```
table(wisc.km$cluster, wisc.hclust.clusters)
```

```
##      wisc.hclust.clusters
##      1  2  3  4  5
## 1  68  3 365  2  0
## 2 109  2  18  0  2
```


Q14. How well does k-means separate the two diagnoses? How does it compare to your hclust results?

It works well, but less accurate than hclust.

## Combining methods

Here we aim to combine our PCA results with clustering. Essentially, we are going to cluster in “PC”, that is cluster on the results `wisc.pr`.

```
plot(wisc.pr$x[,1:2], col= diagnosis)
```




## Clustering on PCA results

I will use 4 PCs and `hclust()` and `dist()` as an input.

```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:4]), method = "ward.D2")
```

```
plot(wisc.pr.hclust)  
abline(h=80, col="red")
```

## Cluster Dendrogram



```
dist(wisc.pr$x[, 1:4])  
hclust (*, "ward.D2")
```

Let's find our cluster membership vector by cutting this tree into k=2 groups.


```
grps <- cutree(wisc.pr.hclust, k=2)  
table(grps)
```

```
## grps  
##    1    2  
## 171 398
```


```
table(grps, diagnosis)
```

```
##      diagnosis  
## grps    B    M  
##    1    6 165  
##    2 351   47
```

```
plot(wisc.pr$x[,1:2], col=grps)
```



```
plot(wisc.pr$x[,1:2], col=diagnosis)
```



To match things up we can turn our groups into a factor and reorder the levels so cluster 2 comes first and thus gets the first color (black) and cluster 1 gets the second color (red).


```
g <- as.factor(grps)
levels(g)
```

```
## [1] "1" "2"
```

```
g <- relevel(g,2)
levels(g)
```


```
## [1] "2" "1"
```

```
plot(wisc.pr$x[,1:2], col=g)
```



```
plot(wisc.pr$x[,1:2], col=diagnosis)
```





diagnosis

grps B M 1 6 165 2 351 47

TP: 165, FP:6 TN: 351, FN: 47

**Accuracy**, essentially how many did we get correct?

**Sensitivity**:  $TP/(TP+FN)$

**Specificity**:  $TN/(TN+FN)$

```
# **Accuracy**: (TP+TN)/total cases
# **Sensitivity**: TP/(TP+FN)
# **Specificity**: TN/(TN+FN)
Accuracy <- (165+351)/nrow(wisc.data)*100
Sensitivity <- (165)/(165+47)*100
Specificity <- (351)/(351+47)*100
evulation <- rbind(Accuracy, Sensitivity, Specificity)
evulation
```

```
##           [,1]
## Accuracy   90.68541
## Sensitivity 77.83019
## Specificity 88.19095
```

Use the distance along the first 7 PCs for clustering i.e. `wisc.pr$x[, 1:7]`

```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:7]), method="ward.D2")
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
table(wisc.pr.hclust.clusters, diagnosis)
```

```
##              diagnosis
## wisc.pr.hclust.clusters  B  M
##              1  28 188
##              2 329  24
```

```
# **Accuracy**: (TP+TN)/total cases
# **Sensitivity**: TP/(TP+FN)
# **Specificity**: TN/(TN+FN)
Accuracy <- (188+329)/nrow(wisc.data)*100
Sensitivity <- (188)/(188+24)*100
Specificity <- (329)/(329+24)*100
evulation <- rbind(Accuracy, Sensitivity, Specificity)
evulation
```

```
##              [,1]
## Accuracy      90.86116
## Sensitivity    88.67925
## Specificity    93.20113
```

Q15. How well does the newly created model with four clusters separate out the two diagnoses?

It looked good, with higher sensitivity and specificity.

Q16. How well do the k-means and hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the `table()` function to compare the output of each model (`wisc.km$cluster` and `wisc.hclust.clusters`) with the vector containing the actual diagnoses.

```
table(wisc.km$cluster, diagnosis)
```

```
##      diagnosis
##           B  M
##      1 356  82
##      2   1 130
```

```
# **Accuracy**: (TP+TN)/total cases
# **Sensitivity**: TP/(TP+FN)
# **Specificity**: TN/(TN+FN)
Accuracy <- (130+356)/nrow(wisc.data)*100
Sensitivity <- (130)/(130+82)*100
Specificity <- (356)/(356+82)*100
evulation <- rbind(Accuracy, Sensitivity, Specificity)
evulation
```

```
##           [,1]
## Accuracy    85.41301
## Sensitivity  61.32075
## Specificity  81.27854
```

```
table(wisc.hclust.clusters, diagnosis)
```

```
##           diagnosis
## wisc.hclust.clusters  B  M
##           1  12 165
##           2   0   5
##           3 343  40
##           4   2   0
##           5   0   2
```

```
# Accuracy: (TP+TN)/total cases
# Sensitivity: TP/(TP+FN)
# Specificity: TN/(TN+FN)
Accuracy <- (165+343)/nrow(wisc.data)*100
Sensitivity <- (165)/(165+40)*100
Specificity <- (343)/(343+40)*100
evulation <- rbind(Accuracy, Sensitivity, Specificity)
evulation
```

```
##           [,1]
## Accuracy    89.27944
## Sensitivity  80.48780
## Specificity  89.55614
```

Q17. Which of your analysis procedures resulted in a clustering model with the best specificity?  
How about sensitivity?

Using the distance along the first 7 PCs for clustering give the best specificity and sensitivity.

## Prediction

We will use the `predict()` function that will take our PCA model from before and new cancer cell data and project

```
#url <- "new_samples.csv"
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

```
##           PC1      PC2      PC3      PC4      PC5      PC6      PC7
## [1,]  2.576616 -3.135913  1.3990492 -0.7631950  2.781648 -0.8150185 -0.3959098
## [2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
##           PC8      PC9      PC10     PC11     PC12     PC13     PC14
## [1,] -0.2307350 0.1029569 -0.9272861 0.3411457  0.375921 0.1610764 1.187882
```

```
## [2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
##          PC15          PC16          PC17          PC18          PC19          PC20
## [1,] 0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
## [2,] 0.1299153 0.1448061 -0.40509706 0.06565549 0.25591230 -0.4289500
##          PC21          PC22          PC23          PC24          PC25          PC26
## [1,] 0.1228233 0.09358453 0.08347651 0.1223396 0.02124121 0.078884581
## [2,] -0.1224776 0.01732146 0.06316631 -0.2338618 -0.20755948 -0.009833238
##          PC27          PC28          PC29          PC30
## [1,] 0.220199544 -0.02946023 -0.015620933 0.005269029
## [2,] -0.001134152 0.09638361 0.002795349 -0.019015820
```

Now add these new samples to our PCA plot

```
plot(wisc.pr$x[,1:2], col=g)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], labels=c(1,2), col="white")
```

